The Challenges

NZZ & Google Cloud Hackathon · 8–9 September 2026 · Opera House, The Village Garden, 18 Liliencron Street, Windhoek

How Teams and Topics Work

You'll find out your team on the day. Teams are put together in advance and balanced for skills, experience and a mix of NZZ and external builders — so every team has a spread of strengths and people to learn from. Teams are announced at the Draw on Tuesday morning.

Your team then picks its challenge. Once you know who you're working with, you choose which of the challenges below to take on. Nothing is assigned to you.

Eight challenges are described below. Each one is a real, documented problem from inside a working newsroom. You don't need media experience to solve them — you need to understand the friction and build something that removes it.

You'll get the full challenge cards on the day, including the business assumptions behind each problem. Part of your job is to pressure-test those assumptions, not just accept them.

The Four Streams

Stream	What it's about
1 · Next-Gen Customer Support	Reducing friction for readers and the people who support them.
2 · Data & Content Personalisation	Making content meet readers where they are — adaptive formats, personalised delivery, smarter discovery.

3 · Internal Productivity & Automation	Removing manual bottlenecks from editorial and operational work.
4 · Open Innovation / Wildcard	Your own idea, built on GCP, solving a real media problem.

The Challenges

1 · Liquid Story Engine

Stream 2 · Data & Content Personalisation

Reporters and desk editors can't turn text-heavy reporting into other formats (video summaries, audio briefs, newsletters, vertical social clips) without it landing on a handful of specialist editors or simply not happening. The hard part isn't the conversion; it's doing it while preserving a distinctive editorial voice. Today more than 60% of high-impact reporting is published in one static text format and nothing else.

Build: a tool that generates multimodal derivatives from an article — storyboards, scripts, modular briefs — that still sound like the publication wrote them.

Consider using: Vertex AI, Gemini API, etc.

2 · Visual Velocity

Stream 3 · Internal Productivity & Automation

Beat journalists have the data and the story, but not an easy way to spot a visual opportunity in their own draft or turn a messy CSV into a chart, timeline or comparison slider. Existing tools demand data cleaning and design decisions that nobody has time for under deadline. Over 80% of data-relevant articles stay as walls of text.

Build: an assistant that flags visualisation opportunities inside a draft and converts raw data into clean, styled interactive graphics in seconds.

Consider: Looker, Vertex AI, BigQuery, etc.

3 · NZZ FlexRead

Stream 2 · Data & Content Personalisation

There's a daily mismatch between how much time a reader has and how long the article is. Two groups feel it most: busy professionals and commuters who need to navigate quickly, and younger readers who arrive from social media expecting something dynamic and hit a dense static page instead. Roughly 800,000 social clicks a year land on articles, and around 90% bounce immediately — the format contradicts what they expected.

Build: an adaptive reading experience — an "essentials in 60 seconds" mode, real-time simplification of dense passages, or another way to flex the format to the moment without losing the substance.

Consider: Gemini API, Vertex AI, etc.

4 · The Market: ImpactMapper

Stream 4 · Open Innovation / Wildcard

Financially literate private investors want to know how a macro shock (a central bank rate decision, a tariff escalation, a geopolitical conflict) affects the sectors and assets they actually hold. They manage those portfolios elsewhere, so today they get general analysis and are left to draw the connections themselves. Many turn to general-purpose AI tools and get confident, hallucinated answers.

The hard constraint: a publisher **cannot give investment advice**. The tool has to be genuinely useful while staying clearly on the right side of that line.

Build: an interactive way to map a macro event onto asset classes and sectors, so a reader can see plausible knock-on effects — informing, never advising.

Consider: BigQuery, Vertex AI, Gemini API, etc.

5 · OSINT Sentinel

Stream 3 · Internal Productivity & Automation

An OSINT team of five and four data journalists verify social media video, satellite imagery and AI-generated pictures against breaking-news deadlines. There's little shared tooling and no consistent verification indicators, so each person improvises their own workflow. It's manual, fragmented and high-stress — and the cost of getting it wrong is the brand's core promise of verified truth.

Build: a verification workspace that pulls the checks into one place — reverse image tracing, geolocation cross-referencing, AI-generation detection — with a human making the final call.

Consider: Vertex AI, Gemini API, Cloud Vision, etc.

6 · DataCatch

Stream 3 · Internal Productivity & Automation

Trend spotting is reactive: editors monitor social media or go on hunches, because connecting live search trends to public datasets needs data-science skills and hours of digging. Stories get published after the peak of search interest, and a backlog of viable trend-matched angles never gets built.

Build: a proactive hypothesis engine that pairs trending signals with open public datasets and surfaces story ideas with the supporting data already attached.

Consider: BigQuery, Gemini API, Cloud Functions, etc.

Scope note: pick one dataset domain — health or economic data works well. Depth beats breadth.

7 · Community Editor

Stream 1 · Next-Gen Customer Support

Two full-time community editors are entirely consumed by reactive moderation — policing spam and rule violations. That leaves zero capacity to do the valuable part: surfacing well-reasoned arguments and genuinely interesting disagreement. Good contributions stay buried, debate quality stays invisible, and a community that could be a premium experience stays flat.

Build: a system that flips moderation from policing to curation — finding and elevating the best reader perspectives rather than only filtering out the worst.

Consider: Vertex AI, Gemini API, etc

8 · Multi-Asset Archive Search

Stream 2 · Data & Content Personalisation

For a Pro subscriber doing deep research, thirty years of audio and video archive is effectively invisible. Ask a research question and you get text extracts at best — high-value multimedia can't be found, cited or played. The alternative is manually scanning hours of podcast episodes for one reference. The archive that premium subscribers pay most for is the part they can't reach.

Build: semantic search across an audio archive that returns the right *moment*, not just the right file — playable, citable, in context.

Consider: Vertex AI Search, Gemini API, Speech-to-Text, etc.

Scope note: narrow to a podcast/audio archive. Video is a stretch goal. Naming this one is part of your deliverable — the working title is a placeholder.

Bring your own

If none of these grabs your team, build something else. Three requirements: it solves a real media problem, it runs on GCP, and you can explain why it matters.

One direction worth exploring — turning an archive from dead storage into a proactive assistant. On-demand briefings built for a specific moment of need: a market briefing before a client meeting, a country guide before a trip.

A few things worth knowing

- Aim to integrate **three or more distinct GCP services** meaningfully.
- The **business case carries the most weight** in judging. A working prototype with a clear answer to "who does this help, and what is it worth?" beats a slick demo without one.
- Each challenge comes with business assumptions attached. **Testing whether they hold up is part of the work** — the strongest pitches will say which numbers they believe and why.
- All work happens in the **provided GCP sandbox environments**.
- Respect **data privacy and anonymization** standards throughout.
- **Final submission is a hard cut-off.** Anything late can't be judged.

Questions

We have a dedicated Google Chat for the hackathon which you should be part of. In this chat, you'll find dedicated channels where you can make inquiries, discuss general topics, and communicate with other participants.

Google Chat: https://mail.google.com/mail/u/0/#chat/space/AAQATeE5Xy4

Emergency

Axel Mukwena

Email: axel.mukwena@nzz.ch
Phone/WhatsApp: +264816035965

Pieter Wiesner

Email: pieter.wiesner@nzz.ch
Phone/WhatsApp: +264813039677